

Reinforcement Learning and Language Model Alignment

VTR UGE 21

DAY 1 · MODULE 1

Foundations of Reinforcement Learning

Learning from interaction: bandits and exploration, Markov decision processes, values and the Bellman equation, and a small world that can be solved exactly.

Prof. Dr. Utku Kose

Full Professor, Department of Computer Engineering, Süleyman Demirel University, Isparta, Türkiye
Founding Director, AI Application and Research Center (YAZEM), Süleyman Demirel University
Head of the Computer Science Division, Department of Computer Engineering, Süleyman Demirel University
Additional affiliations: University of North Dakota (USA), Universidad Panamericana (Mexico City, Mexico),
Vel Tech University (Chennai, India)
IEEE Senior Member, ACM Professional Member

ORCID 0000-0002-9652-6415 · utkukose.com · github.com/utkukose
utkukose@sdu.edu.tr · utku.kose@und.edu · ukose@up.edu.mx · utkukose@gmail.com

Overview

Reinforcement learning is learning from interaction: An agent acts, the environment answers with a new situation and a reward, and the agent improves its behaviour to collect more reward over time [1]. The same idea trains game-playing programs, controls robots and, as Days 4 and 5 show, aligns language models with human preferences. This day introduces the vocabulary of the field on small problems: bandits, where the only question is which action to try [2, 3]; Markov decision processes, where actions also change the situation [1]; and values with the Bellman equation [4], which allow a small world to be solved exactly. That exact solution becomes the yardstick for every method of the week.

Estimated study time: 6 to 8 hours.

Learning outcomes

By the end of the day, students are expected to describe the agent-environment interface and state three ways in which the reward hypothesis can fail, to implement and compare epsilon-greedy, the upper confidence bound (UCB) strategy and Thompson sampling by cumulative regret over many seeds, to formulate sequence generation as a Markov decision process, to write the Bellman expectation and optimality equations and solve small processes with policy and value iteration, to obtain the exact optimum of TokenWorld as a reference for later days, and to report a stochastic experiment with seeds, spread and sample size. In Python, students are expected to use values, variables, lists, loops, decisions, functions and dictionaries to simulate and score a bandit.

Python in this day

The Colab notebook opens with Python step 1: Values, variables, lists and decisions, right after its setup section. It covers values and types, lists, **if** statements, **for** loops, functions and dictionaries, applied to a three-armed bandit played with epsilon-greedy and scored by its regret [5]. After the step, the notebook continues with the full bandit comparison and the two exact solutions. Students who have never programmed should first open the start-here notebook of the course.

How these notes point to the Python code. Every method of this lecture is also written as Python code in the Colab notebook of the day, the file `NB01_foundations_of_rl.ipynb`. A Colab notebook is a document of text and Python code that runs in the web browser, with nothing to install. It opens with the button Open in Colab on the day overview page of the course, <https://github.com/utkukose/rl-alignment-veltech-NB-lecture/tree/main/days/day-01>. The code of the notebook is divided into numbered sections. In these notes, each box titled *Python code in the Colab notebook, Section N* names the section whose code carries out what the text above the box explains, and says what to run and what to look at there. The animations that the notes mention are interactive figures on the lecture page of the day, <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lecture.html>. The interactive lab, with the practice parts and the self-assessment of the day, is at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html>.

Python step 1: Values, variables, lists and decisions

A notebook is a sequence of cells. Text cells explain, and code cells hold Python instructions that run when the cell is executed with Shift+Enter, with the output shown below the cell [7]. This step plays a three-armed bandit with epsilon-greedy and scores it by its regret [5].

Values, lists and variables

```
win_rates = [0.30, 0.45, 0.40] # true chance of a reward per arm, unknown to the agent
pulls = [0, 0, 0] # how often each arm was pulled
wins = [0, 0, 0] # how many rewards each arm gave
estimates = [0.0, 0.0, 0.0] # the agent's estimate of each win rate
print(len(win_rates), "arms; best true win rate:", max(win_rates))
```

Output

```
3 arms; best true win rate: 0.45
```

A list holds values in order between square brackets, and positions start at 0, so `win_rates[1]` is 0.45. `len` returns the length and `max` the largest element. The agent never reads `win_rates`; it only sees rewards.

A decision: explore or exploit

```
import random
rng = random.Random(0) # a reproducible random number generator
eps = 0.1
if rng.random() < eps:
    arm, reason = rng.randrange(3), "explore"
else:
    arm, reason = estimates.index(max(estimates)), "exploit"
print(f"chose arm {arm} to {reason}")
```

Output

```
chose arm 0 to exploit
```

`rng.random()` returns a number between 0 and 1, so the condition is true with probability `eps`. `rng.randrange(3)` returns 0, 1 or 2. `estimates.index(max(estimates))` finds the position of the largest estimate, the arm that looks best. Two values can be assigned at once, separated by a comma.

A loop of 300 pulls

```
for t in range(300):
    if rng.random() < eps:
        arm = rng.randrange(3)
    else:
        arm = estimates.index(max(estimates))
    reward = 1 if rng.random() < win_rates[arm] else 0
    pulls[arm] += 1
    wins[arm] += reward
    estimates[arm] = wins[arm] / pulls[arm]
print("pulls:", pulls)
print("estimates:", [round(e, 3) for e in estimates])
```

Output

```
pulls: [33, 10, 257]
estimates: [0.152, 0.4, 0.447]
```

The loop runs its block 300 times. `1 if rng.random() < win_rates[arm] else 0` is a conditional expression: It gives the reward 1 with the win rate of the arm as probability, and 0 otherwise. `pulls[arm]`

`+= 1` adds one to an element of the list. The estimate of an arm is its share of wins so far, updated after every pull. The square brackets in the last line build a new list of rounded values, a list comprehension.

A function for the regret

```
def regret(pulls, win_rates):
    """Expected reward lost against always pulling the best arm."""
    best = max(win_rates)
    return sum(n * (best - p) for n, p in zip(pulls, win_rates))

print("regret after 300 pulls:", round(regret(pulls, win_rates), 2))
```

Output

```
regret after 300 pulls: 17.8
```

`def` defines a function with parameters and a `return` value, and the text in triple quotation marks documents it. `zip` walks two lists in step. Each pull of an arm that is worse than the best one adds the difference of their win rates to the regret. The function therefore computes the regret of the lecture, grouped by arm: $R_T = \sum_a N_a (\mu^* - \mu_a)$, with `pulls` for the numbers of pulls N_a , `win_rates` for the probabilities μ_a and `best` for μ^* . In this run most pulls went to the arm with a win rate of 0.40 rather than the best arm with 0.45: The noisy early estimates favoured it, and ten percent of exploration did not correct them within 300 pulls.

A dictionary of results

```
summary = {"arms": 3, "epsilon": eps, "total reward": sum(wins),
           "regret": round(regret(pulls, win_rates), 2)}
for key, value in summary.items():
    print(f"{key:13s} {value}")
```

Output

```
arms          3
epsilon       0.1
total reward  124
regret        17.8
```

A dictionary maps keys to values between curly braces, and `items` returns key and value pairs for a loop. In an f-string, `{key:13s}` pads the text to thirteen characters so that the values line up.

Quick check. Why does an epsilon-greedy agent with `eps = 0` risk a large regret?

Learning from interaction

In supervised learning, every example comes with the right answer. In reinforcement learning, abbreviated RL, nobody gives the right action; the agent only receives a reward after acting, and the reward may come much later than the action that earned it [1]. The agent and the environment exchange three things in a loop: The environment shows a state, the agent chooses an action, and the environment returns a reward and the next state. The behaviour of the agent, the rule that maps states to actions, is called its policy, and its goal is to maximise the total reward over time, called the return.

In everyday life: learning from rewards

A child who learns to ride a bicycle gets no instruction for every movement of the handlebar, only the feedback of staying up or falling over. A chess player learns only at the end of a game whether all the moves together won or lost. Both learn from rewards, and in chess the reward comes long after the moves that earned it. This delay is what makes reinforcement learning hard: The learner has to find out which of its earlier actions deserve the credit.

TokenWorld, the project of the course

The course follows one small world through all five days. In TokenWorld an agent writes a short sentence one token at a time, and the finished sentence receives a reward for its quality. Writing a sentence token by token is what a language model does, so TokenWorld is a miniature of the alignment of language models that Days 4 and 5 treat.

Each day takes the project one step further. Day 1 describes TokenWorld as a decision problem and solves it exactly, which gives the optimum against which every later method is scored. Day 2 takes away the knowledge of the rules and learns the values of TokenWorld from experience. Day 3 learns the policy itself. Day 4 enlarges TokenWorld to a grammar of 1512 sentences, lets a small neural language model write them and aligns the model with preferences through a learned reward. Day 5 aligns the same model without a reward model and asks how the result of an alignment is measured.

These words have fixed symbols. Time is counted in steps $t = 0, 1, 2, \dots$. At step t the agent sees the state s_t , chooses the action a_t , and receives the reward r_{t+1} , a number, together with the next state s_{t+1} . The policy is written with the Greek letter pi, $\pi(a | s)$: the probability that the agent chooses action a when it is in state s . The vertical bar reads *given*. The return from step t on, written G_t , is the sum of the rewards that follow:

$$G_t = r_{t+1} + r_{t+2} + r_{t+3} + \dots$$

For example, an agent that receives the rewards 0, 0 and 5 in its next three steps and nothing afterwards has the return $G_t = 0 + 0 + 5 = 5$. Reinforcement learning searches for a policy whose return is as large as possible on average.

You, the agent action: **right** → Environment → next state and reward: **row 1, column 6; +10**

↑
← ↓ →

New episode Random agent Show the map

Step 8: You chose **right**. The robot reached the charger: reward +10. The episode ends with the return $G = 3$, the sum of all its rewards.

episode: 2 step t: 8 last reward: +10 return G : 3 best return: 3

*Interactive scene on the lecture page: The robot in the fog. A robot in a grid of four rows and six columns, steered with the arrow buttons or the arrow keys. Every move earns the reward -1 , the charger $+10$ and a pit -10 , and now and then the floor makes the robot slip in another direction. Fog hides a cell until the robot first enters it, so the map is learned from interaction alone. The readouts give the return of the episode, the sum of its rewards, and the button *Random agent* plays one episode with random actions for comparison.*

Behind this goal stands the reward hypothesis: Every goal can be described as the maximisation of the expected sum of one number, the reward [1]. Three objections to it matter for Days 4 and 5. First, real goals are several at once, such as fast, safe and cheap, and merging them into one number needs weights that somebody has to choose. Second, a reward that is written down can differ from what was really wanted. An agent can then collect reward without doing the intended thing, which is called reward hacking. Third, for a task such as writing a good answer nobody can write the reward as a formula, so it has to be learned from people who compare answers.

Check your understanding (1). What does a reinforcement learning agent receive that tells it how well it acted?

- A. The correct action
- B. A reward, possibly long after the action
- C. A label for each state
- D. The model of the environment

Bandits: exploration against exploitation

The simplest problem has one state and several actions, like a row of slot machines, each paying with an unknown probability: a multi-armed bandit. The name comes from the one-armed bandit, a nickname of the slot machine with its lever. Each machine, and so each action, is called an arm, and playing an arm once is called a pull. There are K arms. Arm a pays a reward of 1 with a probability μ_a , the Greek letter mu, and a reward of 0 otherwise. The agent does not know these probabilities. The best arm has the highest one, written $\mu^* = \max_a \mu_a$.

The agent can learn the probabilities only from what it has seen. Let N_a be the number of times that arm a has been pulled so far and S_a the number of rewards it has paid. The estimate of its probability, marked with a hat, is the share of pulls that paid:

$$\hat{\mu}_a = \frac{S_a}{N_a}$$

→ This formula with numbers, step by step: [Lab Part A](#), step 3.

An arm that paid 6 times in 10 pulls has the estimate $\hat{\mu}_a = 6 / 10 = 0.6$. After a few pulls the estimate can be far from the truth: An arm with a true probability of 0.45 pays only once in its first five pulls in about one case out of five, which gives an estimate of 0.2. The agent must therefore balance exploitation, choosing the arm that looks best now, with exploration, trying other arms that might be better. Four strategies strike this balance in different ways. In all of them a_t is the arm chosen at pull number t .

In everyday life: exploration and exploitation

Choosing a restaurant is a bandit problem. The favourite place is the safe choice, exploitation, and a new place may turn out better or worse, exploration. Someone who only ever visits the favourite never finds out whether the new place is better, and someone who always tries new places rarely enjoys the best one. An online shop that shows two designs of a page to its visitors and keeps the one that sells more faces the same choice.

Greedy. The greedy strategy always pulls the arm with the highest estimate:

$$a_t = \arg \max_a \hat{\mu}_a$$

→ This formula with numbers, step by step: [Lab Part A](#), step 4.

The symbol $\arg \max$ returns the arm at which the estimate is largest, while $\max$ would return the largest estimate itself. A greedy agent never explores and can lock onto a poor arm after a few unlucky early results.

Epsilon-greedy. The epsilon-greedy strategy, also written ε -greedy with the Greek letter epsilon, explores at random a small share of the time:

$$a_t = \begin{cases} \text{an arm chosen at random} & \text{with probability } \varepsilon \\ \arg \max_a \hat{\mu}_a & \text{with probability } 1 - \varepsilon \end{cases}$$

→ This formula with numbers, step by step: [Lab Part A](#), step 5.

With $\varepsilon = 0.1$, one pull in ten is an exploration and nine in ten are greedy. This is the value used in the experiments of the day. With $\varepsilon = 0$ the strategy is greedy, and with $\varepsilon = 1$ it is purely random.

Upper confidence bound. The upper confidence bound strategy, abbreviated UCB, explores where the uncertainty is largest [2]. It adds to each estimate a bonus that is large for arms that have been pulled rarely, and pulls the arm with the highest sum:

$$a_t = \arg \max_a \left[\hat{\mu}_a + c \sqrt{\frac{\ln t}{N_a}} \right]$$

→ This formula with numbers, step by step: [Lab Part A](#), steps 6 to 10.

Here t is the number of the pull that is being decided, 1 for the first pull, $\ln$ is the natural logarithm and c is a constant that sets the amount of exploration. The experiments of the day use $c = 2$ and pull every arm once before the rule applies, so that no N_a is zero. The sum is an optimistic guess of how good the arm could still be: the upper end of a confidence interval around the estimate, which gives the strategy its name. An arm that has been pulled often has a small bonus and is judged by its estimate. The bonus of a neglected arm keeps growing with $\ln t$, so the arm is pulled again sooner or later.

Worked example: greedy against UCB

Three arms have been pulled 20, 10 and 2 times and have paid 9, 6 and 0 times. These are $20 + 10 + 2 = 32$ pulls, so the next pull has the number $t = 33$. The estimates are $9 / 20 = 0.45$, $6 / 10 = 0.60$ and $0 / 2 = 0.00$. Greedy pulls the second arm, the one with the largest estimate.

UCB with $c = 2$ adds the bonus $2 \sqrt{(\ln 33 / N_a)}$, with $\ln 33 = 3.50$. In this and the following worked examples, intermediate values are rounded to two or three decimals, and each step continues with the rounded value. For the first arm, $3.50 / 20 = 0.175$, $\sqrt{(0.175)} = 0.418$ and the bonus is $2 \times 0.418 = 0.84$. For the second arm, $3.50 / 10 = 0.35$, $\sqrt{(0.35)} = 0.592$ and the bonus is $2 \times 0.592 = 1.18$. For the third arm, $3.50 / 2 = 1.75$, $\sqrt{(1.75)} = 1.323$ and the bonus is $2 \times 1.323 = 2.65$.

The sums of estimate and bonus are $0.45 + 0.84 = 1.29$, $0.60 + 1.18 = 1.78$ and $0.00 + 2.65 = 2.65$. UCB pulls the third arm: Two pulls are too little evidence to give up on it.

Interactive lab, Part A.

The same calculation, together with the choice of epsilon-greedy, with sliders for the pulls and the rewards of each arm, for ϵ and for the constant c . Each estimate, each bonus and each sum is shown one step at a time, and a figure stacks the bonus of every arm on its estimate. Part A opens at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html#pa>.

Thompson sampling. Thompson sampling keeps for every arm a belief about its unknown probability, in the form of a probability distribution [3]. For rewards of 0 and 1 the suitable family is the Beta distribution, a distribution over the numbers between 0 and 1 with two parameters. The agent starts with a uniform belief, and after S_a paid pulls and $N_a - S_a$ unpaid ones its belief about arm a is

$$\mu_a \sim \text{Beta}(1 + S_a, 1 + N_a - S_a)$$

The sign $\sim$ reads *is distributed as*. The mean of this belief, $(1 + S_a) / (2 + N_a)$, lies close to the estimate, and the belief becomes narrower with every pull. At each pull the agent draws one random number, written $\mu_a^\sim$ with a tilde, from the belief of every arm and pulls the arm with the largest draw, $a_t = \arg \max_a \mu_a^\sim$. An arm with few pulls has a wide belief, so its draw is sometimes the largest and the arm is explored. As a result, each arm is pulled with the probability that it is the best one.

A belief $\text{Beta}(\alpha, \beta)$, with the Greek letters alpha and beta for its two parameters, has the mean and the standard deviation, a measure of its width,

$$\text{mean} = \frac{\alpha}{\alpha + \beta}, \quad \text{standard deviation} = \sqrt{\frac{\alpha \beta}{(\alpha + \beta)^2 (\alpha + \beta + 1)}}$$

Worked example: two beliefs in Thompson sampling

One arm has paid 6 times in 10 pulls. Its belief is $\text{Beta}(1 + 6, 1 + 4) = \text{Beta}(7, 5)$, with the mean $7 / (7 + 5) = 7 / 12 = 0.58$. For its standard deviation, $\alpha \beta = 7 \times 5 = 35$ and $(\alpha + \beta)^2 (\alpha + \beta + 1) = 12^2 \times 13 = 1872$, so the standard deviation is $\sqrt{(35 / 1872)} = \sqrt{(0.0187)} = 0.14$. A second arm has paid once in 2 pulls. Its belief is $\text{Beta}(1 + 1, 1 + 1) = \text{Beta}(2, 2)$, with the mean $2 / 4 = 0.50$ and the standard deviation $\sqrt{(4 / (4^2 \times 5))} = \sqrt{(0.05)} = 0.22$.

The first arm looks better, but the belief about the second arm is much wider. The draw of the second arm is the larger one with a probability of 0.38, in about 38 of 100 pulls, so this arm is still pulled often. This probability is an integral over the two beliefs and is computed numerically.

After ten times as many pulls with the same shares, 60 rewards in 100 pulls and 10 in 20, the beliefs are $\text{Beta}(61, 41)$ and $\text{Beta}(11, 11)$. Their standard deviations are $\sqrt{(2501 / (102^2 \times 103))} = 0.05$ and $\sqrt{(121 / (22^2 \times 23))} = 0.10$. The second arm now wins the draw in only about 20 of 100 pulls: Exploration fades as the evidence grows.

Regret. The price of learning is measured by regret, the reward lost compared with always pulling the best arm. One pull of arm a_t loses $\mu^* - \mu_{a_t}$ on average, and the cumulative regret after T pulls is the sum of these losses:

$$R_T = \sum_{t=1}^T (\mu^* - \mu_{a_t})$$

The sign Σ means: Add the term for $t = 1, 2, \dots, T$. For example, when the best arm pays with probability 0.45, then 33 pulls of an arm with 0.30 and 257 pulls of an arm with 0.40 cost $33 \times 0.15 + 257 \times 0.05 = 4.95 + 12.85 = 17.8$. The regret curve of a strategy that has found the best arm becomes flat, and the curve of a strategy that keeps pulling a worse arm rises in a straight line.

Interactive lab, Part B.

A bandit game against the strategies of this section: five machines with hidden probabilities, fifty pulls for each player, and a scoreboard with the reward and the regret of both. Part B opens at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html#pb>.

The experiment in detail. The four strategies are compared on a bandit with $K = 6$ arms whose probabilities lie close together: 0.30, 0.35, 0.38, 0.40, 0.42 and 0.45. The best arm is the sixth, with $\mu^* = 0.45$. One run consists of $T = 1500$ pulls. At every pull the experiment does four things. It chooses an arm by the rule of the strategy. It draws the reward, 1 with the probability of that arm and 0 otherwise. It adds one to N_a and the reward to S_a . It records the loss $\mu^* - \mu_{a_t}$ of the pull. The cumulative regret after each pull is the running sum of these losses.

Each strategy is run thirty times. A run is identified by its seed, the number that starts the random number generator: The same seed always produces the same run, and the seeds 0 to 29 give thirty different runs. For the final regrets $x_1, \dots, x_n$ of the $n = 30$ runs, the experiment reports the mean, written $\bar{x}$ with a bar, the standard deviation (SD), the usual measure of spread, and the largest value, the worst seed:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \quad \text{SD} = \sqrt{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$$

The 10th and the 90th percentile, which bound the bands of the figure, are the values below which 10 and 90 percent of the runs fall.

Worked example: the first pulls of UCB in this experiment

In the run with seed 0, UCB first pulls every arm once, in order. Arms 2, 3 and 4 pay, and arms 1, 5 and 6 do not. The losses of these six pulls are $0.15 + 0.10 + 0.07 + 0.05 + 0.03 + 0 = 0.40$, the regret so far.

Pull 7 is the first that uses the rule. Every arm has $N_a = 1$, so every bonus is $2 \sqrt{(\ln 7 / 1)}$. With $\ln 7 = 1.946$ and $\sqrt{1.946} = 1.395$, the bonus is $2 \times 1.395 = 2.79$. The sums are $1 + 2.79 = 3.79$ for the three arms that paid and $0 + 2.79 = 2.79$ for the others. Among equal sums the first arm is taken, here arm 2, and this pull does not pay.

At pull 8, arm 2 has the estimate $1 / 2 = 0.5$ and, after two pulls, a smaller bonus. With $\ln 8 = 2.079$, its bonus is $2 \sqrt{(2.079 / 2)} = 2 \sqrt{1.04} = 2 \times 1.020 = 2.04$, a sum of $0.5 + 2.04 = 2.54$. Arms 3 and 4 have the bonus $2 \sqrt{(2.079)} = 2 \times 1.442 = 2.88$ and the sum $1 + 2.88 = 3.88$, and UCB pulls arm 3, the first of the two. With bonuses this large, UCB returns to every arm many times before the estimates decide.

Animation on the lecture page: Mean cumulative regret of three strategies over 1500 pulls and twenty seeds. Change ϵ and the constant c of UCB and watch the ranking.




Figure. Cumulative regret of four strategies on six close arms, with the band between the 10th and 90th percentile over thirty seeds.

Strategy	Mean final regret	Standard deviation	Worst seed
Greedy	51.5	50.0	223.5
Epsilon-greedy, $\epsilon = 0.1$	37.9	27.1	138.1
UCB, $c = 2$	81.7	3.9	87.8
Thompson sampling	45.9	11.3	72.2

The figure shows the mean cumulative regret of each strategy, with a band between the 10th and the 90th percentile of the thirty runs, and the table gives the regret after the last pull. For comparison, pulling arms

purely at random loses $0.45 - 0.383 = 0.067$ per pull on average, about 100 over 1500 pulls. Epsilon-greedy ends with the lowest mean regret, about 38, followed by Thompson sampling with 46, greedy with 52 and UCB with 82. Greedy varies most: In its best seed it finds the best arm at once and loses less than 1, and in its worst seed it locks onto a poor arm and loses 224. UCB with $c = 2$ varies least, but on arms this close and with so few pulls its bonus buys more exploration than it repays. Its guarantee concerns the long run [2].

The constant c decides how much UCB explores. The same experiment with other values of c gives:

Constant c	Mean final regret	Worst seed
0.25	42.8	112.7
0.5	40.6	60.3
1	63.1	89.0
2	81.7	87.8
4	89.8	94.2

A small constant explores too little and behaves like greedy in its bad seeds. A large one explores for too long. Here $c = 0.5$ is the best compromise, with a mean regret of 41 and a worst seed of only 60. A mean alone would hide such differences. For this reason the standard deviation and the worst seed are reported next to every mean.

Python code in the Colab notebook, Section 1.

Open Section 1 of the Colab notebook and run it. The function `run_bandit` contains the four rules, each as one line of Python, and the cell prints the two tables and draws the figure shown here. Then change `c=2.0` or `eps=0.1` in the first line of the function and run the cell again to see the ranking change.

Check your understanding (2). Why can a purely greedy strategy end with a large regret?

- A. It explores too much
- B. It can lock onto a poor arm after a few unlucky early results
- C. It is too slow
- D. It needs a model

Check your understanding (3). An arm has been pulled 4 times and has paid 3 times. At pull number $t = 100$ and with $c = 2$, what is its upper confidence bound?

- A. 0.75
- B. About 1.82
- C. About 2.90
- D. About 5.36

Markov decision processes

In most problems an action also changes the situation. A Markov decision process, abbreviated MDP, describes such a problem with five ingredients [1]: a set of states, a set of actions, transition probabilities,

rewards and a discount factor. The transition probability $P(s' | s, a)$ is the probability that the next state is s' when the agent takes action a in state s . The reward $r(s, a, s')$ is the number that the agent receives for this step. The Markov property says that the next state depends only on the current state and action, not on the whole history:

$$P(s_{t+1} | s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = P(s_{t+1} | s_t, a_t)$$

The discount factor γ , the Greek letter gamma, is a number between 0 and 1 that makes rewards count less the later they come. With it the return becomes a discounted sum:

$$G_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$$

A reward that arrives k steps later is multiplied by γ^k . One run from a start state to an end state is called an episode.

In everyday life: the Markov property and discounting

A chess position has the Markov property: The possible next positions depend only on the position on the board and the move, not on the moves that led there. Discounting is familiar from money: 100 rupees today are usually preferred to 100 rupees in a year. A discount factor of 0.9 per step makes a reward one step later worth 90 percent of the same reward now, and a reward ten steps later only $0.9^{10} = 0.35$ of it.

Worked example: a discounted return

An episode gives the rewards 2, 0 and 10 and then ends. With $\gamma = 0.9$ the return is $2 + 0.9 \times 0 + 0.9^2 \times 10$. Here $0.9^2 = 0.81$ and $0.81 \times 10 = 8.1$, so the return is $2 + 0 + 8.1 = 10.1$. With $\gamma = 0.5$ the same rewards give $2 + 0.5 \times 0 + 0.5^2 \times 10 = 2 + 0 + 0.25 \times 10 = 4.5$: The late reward counts much less. With $\gamma = 1$ nothing is discounted, and the return is the plain sum, 12.

The project of the course, TokenWorld, is a Markov decision process of this kind: The agent writes a short sentence one word at a time and is rewarded at the end for its quality. It is a miniature of how a language model is trained with reinforcement learning on Day 4. A unit of text, here a word, is called a token. The vocabulary has four tokens: **the**, **cat**, **sat** and **<eos>**, the end-of-sentence token. The state is the sentence written so far, the action is the next token, and the transition is certain: The token is appended. An episode ends when the agent writes **<eos>** or when the sentence has four tokens. Only then does the reward arrive, computed by the rules of the table, and no reward is discounted, $\gamma = 1$. TokenWorld has 161 states: the empty sentence and 4, 12, 36 and 108 sentences of one, two, three and four tokens. Of these, 121 are finished sentences.

One episode of TokenWorld: the best sentence, written token by token




Figure. One episode of TokenWorld along the best sentence. Each state is the sentence written so far, each action appends one token, and the only reward arrives when the sentence is finished. The lower row counts all sentences of each length.

Property of the finished sentence	Reward
It starts with the	+1.0, otherwise -0.5
cat appears exactly once	+1.5
cat does not appear	-0.5
cat appears more than once	-0.8 for every repetition
sat appears, and the first cat stands before it	+1.0
Each word, the end token not counted	-0.25
The sentence has no word at all	-1.0 as the whole reward

For example, the sentence **the cat sat <eos>** earns $1.0 + 1.5 + 1.0 - 3 \times 0.25 = 2.75$, and the sentence **cat cat <eos>** earns $-0.5 - 0.8 - 2 \times 0.25 = -1.8$.

The count can be followed step by step. After the first token there are 4 sentences, of which 3 are unfinished, because one of them is **<eos>** alone. Each unfinished sentence can be continued with 4 tokens. This gives $3 \times 4 = 12$ sentences of two tokens, 9 of them unfinished, then $9 \times 4 = 36$ of three tokens, 27 of them unfinished, and $27 \times 4 = 108$ of four tokens. Together with the empty sentence these are $1 + 4 + 12 + 36 + 108 = 161$ states. Finished are the $1 + 3 + 9 = 13$ shorter sentences that end with **<eos>** and all 108 sentences of four tokens, in total 121.

Python code in the Colab notebook, Section 2.

Open Section 2 of the Colab notebook and run it. The function **score_sequence** is the reward table above written as Python code, and the function **step** appends one token to the sentence. The second cell counts the 161 states and prints the five best and the three worst sentences with their rewards.

Check your understanding (4). What does the discount factor control?

- A. The number of actions
- B. How much distant rewards count compared with near ones
- C. The learning rate
- D. The exploration

Values and the Bellman equation

The value of a state is the return the agent can expect from it when following its policy. As a formula, with E_π for the average over everything that can happen while the agent follows the policy π :

$$V^\pi(s) = \mathbb{E}_\pi [G_t \mid s_t = s]$$

Interactive lab, Part C.

Six descriptions from the life of a delivery robot, each to be matched with its ingredient of a decision problem: state, action, reward, policy, value or discount factor. Part C opens at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html#pc>.

In everyday life: values

A navigation app shows the remaining travel time from every junction of a route. The time from one junction is the time of the road to the next junction plus the remaining time from there. This is the Bellman equation, with travel times in place of rewards. The app can compute all these times backwards from the destination, which is the idea of value iteration.

Bellman observed that values satisfy a recursive relation: The value of a state equals the expected reward of the next step plus the discounted value of the state that follows [4]. For a given policy π this relation is the Bellman expectation equation:

$$V^\pi(s) = \sum_a \pi(a \mid s) \sum_{s'} P(s' \mid s, a) [r(s, a, s') + \gamma V^\pi(s')]$$

The two sums average over the action that the policy chooses and over the next state that follows. For the best possible policy, whose values are written V^* , the average over the actions is replaced by the best action, which gives the Bellman optimality equation:

$$V^*(s) = \max_a \sum_{s'} P(s' \mid s, a) [r(s, a, s') + \gamma V^*(s')]$$

The formula is read from the inside. For one action a , every possible next state s' contributes its reward plus its discounted value, weighted by its probability. The sum over s' is therefore the expected outcome of the action, called its action value $Q^*(s, a)$. The max over a then picks the best action, and the best policy π^* takes that action:

$$Q^*(s, a) = \sum_{s'} P(s' \mid s, a) [r(s, a, s') + \gamma V^*(s')]$$

$$V^*(s) = \max_a Q^*(s, a)$$

$$\pi^*(s) = \arg \max_a Q^*(s, a)$$

Worked example: one Bellman backup

One application of the right-hand side to one state is called a backup. A state has two actions, every step costs 1, and $\gamma = 0.9$. The action *left* leads with certainty to a state of value 4, so its value is $-1 + 0.9 \times 4 = -1 + 3.6 = 2.6$. The action *right* leads with probability 0.7 to a state of value 10 and with probability 0.3 to a state of value -5. The first outcome is worth $-1 + 0.9 \times 10 = -1 + 9 = 8$ and the second $-1 + 0.9 \times (-5) = -1 - 4.5 = -5.5$. Weighted by their probabilities, the value of *right* is $0.7 \times 8 + 0.3 \times (-5.5) = 5.6 - 1.65 = 3.95$.

The value of the state is the larger of the two, 3.95, and the best action is *right*.

Interactive lab, Part D.

Two calculations by hand with the formulas of this section, a discounted return and a Bellman backup. The lab checks each value. Part D opens at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html#pd>.

When the transitions are known, the equation can be solved by repeating it until the values stop changing, which is value iteration, a form of dynamic programming. All values start at zero. One sweep visits every state and replaces its value by the backup computed from the current values:

$$V_{k+1}(s) = \max_a \sum_{s'} P(s' | s, a) [r(s, a, s') + \gamma V_k(s')]$$

→ This formula with numbers, step by step: [Lab Part E](#), steps 1 to 14, one cell of one sweep in each.

The index k counts the sweeps. They stop when the largest change of a value in one sweep, $\max_s |V_{k+1}(s) - V_k(s)|$, falls below a small tolerance. Policy iteration is a second method for the same task. It reaches the same values along another route: It computes the values of the current policy by repeating the expectation equation, switches every state to the action with the highest action value, and repeats both steps until the policy no longer changes.

Worked example: value iteration by hand

A corridor has two cells, A and B, and a goal behind B. The action *forward* moves the agent one cell on with probability 0.8 and leaves it where it is with probability 0.2. Every step has a reward of -1, except the step that reaches the goal, which has a reward of 0, and $\gamma = 0.9$. The goal has the value 0, and both cells start with the value 0. Each sweep computes the two new values from the values of the previous sweep, rounded to three decimals.

Sweep 1 uses the starting values. For B: $0.8 \times (0 + 0.9 \times 0) + 0.2 \times (-1 + 0.9 \times 0) = 0 - 0.2 = -0.2$. For A: $0.8 \times (-1 + 0.9 \times 0) + 0.2 \times (-1 + 0.9 \times 0) = -0.8 - 0.2 = -1$.

Sweep 2 uses the values of sweep 1, -0.2 for B and -1 for A. For B: $0.8 \times 0 + 0.2 \times (-1 + 0.9 \times (-0.2)) = 0.2 \times (-1.18) = -0.236$. For A: $0.8 \times (-1 + 0.9 \times (-0.2)) + 0.2 \times (-1 + 0.9 \times (-1))$, which is $0.8 \times (-1.18) + 0.2 \times (-1.9) = -0.944 - 0.380 = -1.324$.

Sweep 3 uses -0.236 for B and -1.324 for A. For B: $0.2 \times (-1 + 0.9 \times (-0.236)) = 0.2 \times (-1.212) = -0.242$. For A: $0.8 \times (-1.212) + 0.2 \times (-1 + 0.9 \times (-1.324))$, which is $-0.970 + 0.2 \times (-2.192) = -0.970 - 0.438 = -1.408$. The table below continues in the same way until no value changes any more. A second action, *wait*, would have the value $-1 + 0.9 V(s)$, which is lower in every sweep. In sweep 2, for example, it is $-1 + 0.9 \times (-0.2) = -1.18$ for B, less than -0.236, so the max chooses *forward*.

Sweep	$V(B)$	$V(A)$	Largest change of the sweep
1	-0.200	-1.000	1.000
2	-0.236	-1.324	0.324
3	-0.242	-1.408	0.084
4	-0.244	-1.427	0.019
5	-0.244	-1.433	0.006
6	-0.244	-1.434	0.001
7	-0.244	-1.434	0.000

The largest change shrinks by a factor of three to four from one sweep to the next. With a tolerance of 0.001, a change below 0.001, the sweeps stop after sweep 7, at $V(A) = -1.434$ and $V(B) = -0.244$. Without rounding, the values approach -1.4337 and -0.2439 , the exact solution of the two Bellman equations.

Interactive lab, Part E.

Value iteration on this corridor, sweep by sweep, with sliders for the probability of moving forward, the discount factor and the reward of a step. A figure follows the two values until they settle. Part E opens at <https://utkukose.github.io/rl-alignment-veltech-NB-lecture/days/day-01/lab.html#pe>.

The Windy Cliff is a grid of 4 rows and 8 columns. The agent starts in the bottom left corner and has to reach the goal in the bottom right corner, and the cells between them are a cliff. The four actions are up, down, left and right. Every step has a reward of -1 , except the step that reaches the goal, which has a reward of 0 . A step into the cliff has a reward of -100 and puts the agent back at the start. The wind is described by the slip probability, 0.15 in this example: With this probability the wind replaces the chosen action by a push to the left or to the right, each with half of it, 0.075 . The discount factor is 0.99 . The optimal policy walks along the row next to the cliff, the shortest route: A sideways push cannot move the agent down into the cliff from that row. The wind changes the values and not the route. The value of the start falls from -7.7 without wind to -17.9 with a slip probability of 0.15 , because at the start a push to the right ends in the cliff. Value iteration computes the values of all 32 cells with the update above and stops when no value changes by more than 10^{-9} , here after 26 sweeps. Policy iteration ends with the same values after 4 changes of the policy.

Worked example: the two values of the start

Without wind the shortest route has nine steps: one up, seven to the right and one down into the goal. The first eight have a reward of -1 and the last one a reward of 0 . The value of the start is their discounted sum, $-(1 + 0.99 + 0.99^2 + \dots + 0.99^7)$. A sum of this kind, a geometric sum, equals $(1 - 0.99^8) / (1 - 0.99)$. With $0.99^8 = 0.9227$ this is $0.0773 / 0.01 = 7.73$, so the value of the start is -7.73 .

With a slip probability of 0.15 , the action *up* at the start has three outcomes. With probability 0.85 the agent moves up, to a cell whose value is -8.09 . With probability 0.075 the wind pushes it to the right, into the cliff: a reward of -100 and a return to the start. With probability 0.075 the wind pushes it to the left, against the edge of the grid: a reward of -1 , and the agent stays at the start.

The Bellman equation of the start, with V for its own value, is therefore $V = 0.85 \times (-1 + 0.99 \times (-8.09)) + 0.075 \times (-100 + 0.99 V) + 0.075 \times (-1 + 0.99 V)$. The first term is $0.85 \times (-1 - 8.009) = 0.85 \times (-9.009) = -7.658$. The second is $-7.5 + 0.07425 V$, and the third is $-0.075 + 0.07425 V$. Collecting the terms gives $V = -15.233 + 0.1485 V$. Subtracting $0.1485 V$ from both sides gives $0.8515 V = -15.233$, so $V = -15.233 / 0.8515 = -17.89$. The risk of the cliff alone accounts for $0.075 \times 100 / 0.8515 = 7.5 / 0.8515 = 8.8$ of this value.

Animation on the lecture page: Value iteration on the Windy Cliff, sweep by sweep. Change the slip probability and run again: The arrows of the greedy policy keep their route, and the value of the start falls as the wind grows.

Windy Cliff: optimal policy and state values (S start, G goal, X cliff)




Figure. The optimal policy and the state values of the Windy Cliff with a slip probability of 0.15 .

Python code in the Colab notebook, Section 3.

Open Section 3 of the Colab notebook and run its cells. The function `grid_transitions` builds the transition probabilities, and the loop under the comment `value iteration` is the update above. The cells print the value of the start, -17.89 , and draw the figure shown here. Then change `slip=0.15` to `slip=0.0` where `grid_transitions` is called and run the cells again: The route stays, and the value of the start becomes -7.73 .

Check your understanding (5). With a slip probability of 0.15, the start of the Windy Cliff has a value of about minus 18 and the cell above it about minus 8. Why is the start so much worse?

- A. It is one step further from the goal
- B. At the start, a push to the right moves the agent into the cliff, which costs 100 and returns it to the start
- C. Value iteration has not converged there
- D. The discount factor is smaller at the start

TokenWorld, stage 1: an exact optimum as a yardstick

The methods of this day now come back to the project. TokenWorld is a Markov decision process with known transitions and rewards, so value iteration can solve it exactly, and this solution is the first result of the project. The best sentence is **the cat sat <eos>** with a return of 2.75, and value iteration finds the same number as the value of the empty sentence, the state in which every episode starts. The best policy writes this sentence token by token.

TokenWorld, stage 1 of 5




Figure. The stage of the project reached today. Every later method is scored against the optimum of 2.75 found here.

The table shows how the best policy decides. For the sentence written so far, each column gives the action value Q^* of writing that token next and continuing in the best way. The best policy takes the largest value of each row.

Sentence so far	the	cat	sat	<eos>	Best next token
(empty)	2.75	1.50	0.50	-1.00	the
the	2.50	2.75	1.75	0.25	cat
the cat	2.50	0.20	2.75	2.00	sat
the cat sat	2.50	0.20	2.50	2.75	<eos>

Worked example: the last two rows of the table

The values are computed from the end of the sentence backwards. After **the cat sat** every action ends the episode, so each action value is the reward of the finished sentence. Writing **<eos>** gives $1.0 + 1.5 + 1.0 - 3 \times 0.25 = 2.75$. Writing **the** or **sat** gives a sentence of four words with one **cat**, $1.0 + 1.5 + 1.0 - 4 \times 0.25 = 2.50$. Writing **cat** gives **the cat sat cat**, in which **cat** is repeated: $1.0 - 0.8 + 1.0 - 4 \times 0.25 = 0.20$. The value of the state **the cat sat** is the largest of these, 2.75.

One step earlier, after **the cat**, the action **sat** leads to that state, so its action value is 2.75. The action **<eos>** ends with **the cat <eos>**, which earns $1.0 + 1.5 - 2 \times 0.25 = 2.00$. The other two actions lead to states whose values, 2.50 and 0.20, are found in the same way. The largest value is again 2.75, and the best token is **sat**.

A policy that picks each token at random, with probability $1/4$, writes the best sentence with probability $(1/4)^4 = 1/256$, less than once in two hundred tries. Its expected return is the reward of every finished sentence x , weighted by the probability of writing it:

$$\mathbb{E}[G] = \sum_x \left(\frac{1}{4}\right)^{|x|} R(x) = -0.49$$

Here $|x|$ is the number of tokens of the sentence, the end token included, and $R(x)$ is its reward.

The sum has 121 terms and can be followed by the length of the sentences. The single sentence of one token, `<eos>` alone, has the probability $1/4$ and the reward -1 , a contribution of -0.250 . The 3 finished sentences of two tokens have the probability $1/16$ each and rewards that add up to -0.25 , a contribution of -0.016 . The 9 of three tokens, with $1/64$ each, contribute -0.005 . The 108 of four tokens, with $1/256$ each and rewards that add up to -57.05 , contribute -0.223 . The total is -0.493 .

A simulation estimates the same number: 4000 episodes played with the random policy have a mean return of -0.526 . The difference from the exact value is due to chance. One random episode has a standard deviation of 1.20, so the mean of 4000 episodes is uncertain by about $1.20 / \sqrt{4000} = 1.20 / 63.2 = 0.019$. This quantity, the standard deviation divided by the square root of the number of episodes, is called the standard error of the mean. Of the 4000 episodes, 0.35 percent reached the optimum, close to the exact share of $1 / 256$, which is 0.39 percent.

Every learning method of the week is measured against the exact optimum of 2.75, which turns a vague claim that a method works into a number.




Figure. Returns of 4000 rollouts of the random policy, with its mean and the exact optimum.

Python code in the Colab notebook, Section 4.

Open Section 4 of the Colab notebook and run it. The function `value_iteration_tokenworld` computes the action values, and the cell prints the table above row by row. The next cell plays the 4000 random episodes, prints their mean return and draws the histogram shown here.

Your field

Bandits appear wherever options are tested on the fly: doses in a clinical trial, subject lines of an email campaign, tariffs offered to drivers, irrigation schedules or hint styles in a tutoring system. The application of the day poses such a problem for five fields. For medicine, the four arms are four doses to which a patient responds with the probabilities 0.20, 0.35, 0.50 and 0.45. The four strategies of the day

run on these arms for 1000 pulls and over 20 seeds. Epsilon-greedy and Thompson sampling end with mean regrets of about 25 and 26, greedy with 58 and UCB with $c = 2$ with 69.

Python code in the Colab notebook, Section 5.

Open Section 5 of the Colab notebook. Its first line is a switch, `FIELD = "medicine"`. Replace the word by `marketing`, `energy`, `agriculture` or `education` and run the section. The printed table shows which strategy has the lowest mean regret in that field and which has the smallest spread.

Going further (optional)

Results in reinforcement learning vary strongly from one random seed to the next, so a single run proves little [6]. The rest of the week therefore follows three habits: Every method is run over many seeds, the spread is reported next to the mean, and an estimate is checked against an exact value wherever one exists.

TokenWorld puts these habits to work. The best policy writes a random token with probability ε , and each such policy is evaluated with 15 seeds of 400 episodes each:

Share of random tokens ε	Mean return over 15 seeds	Standard deviation of the seeds
0	2.750	0.000
0.1	2.320	0.035
0.3	1.534	0.050
0.6	0.606	0.053
1	-0.478	0.052

With $\varepsilon = 1$ the policy is the random policy, whose exact expected return is -0.493 . The measured -0.478 agrees with it within the uncertainty. One seed averages 400 episodes, so its mean varies from seed to seed by about $1.20 / \sqrt{(400)} = 1.20 / 20 = 0.060$, close to the measured 0.052. The mean of 15 seeds is uncertain by about $0.060 / \sqrt{(15)} = 0.060 / 3.87 = 0.016$, and the two numbers differ by $0.493 - 0.478 = 0.015$.

Python code in the Colab notebook, Section 6.

Open Section 6 of the Colab notebook and run it. The function `study` evaluates one policy over 15 seeds and returns the mean, the standard deviation and the best and the worst seed, and the cell prints the table above. The last cell shows three sliders, for ε , for c and for the number of seeds: Move them and watch the regret curves of the bandit change.

Answers to the knowledge checks

Quick check. It never explores: If the first pulls make a poor arm look best, it keeps pulling that arm and never corrects the estimate of the others.

1. **B.** Rewards are the only feedback, and they may be delayed.
2. **B.** Without exploration, an early wrong impression is never corrected.

3. **C.** The estimate is $3 / 4 = 0.75$. For the bonus, $\ln 100 = 4.605$, $4.605 / 4 = 1.151$, $\sqrt{(1.151)} = 1.073$ and $2 \times 1.073 = 2.15$. The sum is $0.75 + 2.15 = 2.90$.

4. **B.** A discount below one makes rewards count less the later they come.

5. **B.** One more step explains a difference of about one. The rest is the risk of the start: Each attempt to leave it ends in the cliff with a probability of 7.5 percent.

References

- [1] Sutton, R. S., & Barto, A. G. (2018). Reinforcement Learning: An Introduction (2nd ed.). MIT Press.
- [2] Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47(2-3), 235-256.
- [3] Thompson, W. R. (1933). On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25(3-4), 285-294.
- [4] Bellman, R. (1957). A Markovian decision process. *Journal of Mathematics and Mechanics*, 6(5), 679-684.
- [5] Van Rossum, G., & Drake, F. L. (2009). Python 3 Reference Manual. CreateSpace.
- [6] Henderson, P., Islam, R., Bachman, P., Pineau, J., Precup, D., & Meger, D. (2018). Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1), 3207-3214.
- [7] Kluyver, T., Ragan-Kelley, B., Pérez, F., Granger, B., Bussonnier, M., Frederic, J., Kelley, K., Hamrick, J., Grout, J., Corlay, S., et al. (2016). Jupyter Notebooks: A publishing format for reproducible computational workflows. In *Positioning and Power in Academic Publishing: Players, Agents and Agendas* (pp. 87-90). IOS Press.

Prof. Dr. Utku Kose

Full Professor, Department of Computer Engineering, Süleyman Demirel University, Isparta, Türkiye

Founding Director, AI Application and Research Center (YAZEM), Süleyman Demirel University

Head of the Computer Science Division, Department of Computer Engineering, Süleyman Demirel University

Additional affiliations: University of North Dakota (USA), Universidad Panamericana (Mexico City, Mexico), Vel Tech University (Chennai, India)

IEEE Senior Member, ACM Professional Member

ORCID 0000-0002-9652-6415 · utkukose.com · github.com/utkukose

utkukose@sdu.edu.tr · utku.kose@und.edu · ukose@up.edu.mx · utkukose@gmail.com